

# GpuGen: Bringing the Power of GPUs into the Haskell World

Sean Lee

`seanl@cse.unsw.edu.au`

Programming Languages and Systems  
School of Computer Science and Engineering  
The University of New South Wales  
*In affiliation with NVIDIA Corporation*

02 Sep 2008

# Graphics Processing Units (GPUs)

## GPUs

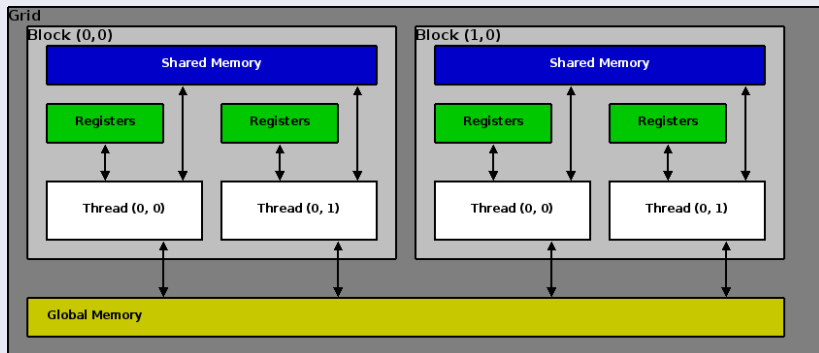
- Superior performance / cost ratio compared to CPUs
- Programmable in SIMD or MIMD style, thus suitable for data parallelism

# Compute Unified Device Architecture (CUDA)

- Available on GeForce 8 (G80) and beyond
- A C++ variant designed for *General Purpose GPU computing* (GPGPU)
- Hardware-based program decomposition:  
Functions to run on GPUs (kernels) are prefixed with `__global__` or `__device__`.
- Specialised memory hierarchy and thread hierarchy

# Compute Unified Device Architecture (CUDA)

## Structure



# Compute Unified Device Architecture (CUDA)

## Block

- No synchronization across multiple blocks
- Limited number of registers per block
- `blockIdx`

## Warp

- A group of threads that run physically in parallel
- `warpSize`

## Thread

- Synchronisation within a block
- Shared memory access within a block
- `threadIdx`, `blockDim`

# Global Memory and Memory Coalescing

## Global Memory

- Global memory access is one of the heaviest overhead in CUDA (400 - 600 cycles)
- Global memory accesses within a half-warp must be coalesced to minimise the number of memory transactions.
- The difference could be an order of magnitude or more.

## Memory Coalescing Rules

- The start address of the block has to be aligned to the multiple of region size.
- The  $k$ th thread in the half-warp must access  $k$ th element in the block that is being read, or not participate in the memory access.

# Shared Memory and Bank Conflicts

## Shared Memory

- The shared memory access is about hundreds of times faster than global memory access.
- The shared memory could be used as the software cache to minimise the global memory accesses or to coalesce the global memory accesses.
- The shared memory is split into equally-sized banks.

## Bank Conflicts Rules

- If more than one threads try to access the elements in the same bank, a bank conflict occurs and those accesses are serialised.
- If all threads in a half-warp read from an address within the same 32-bit word, there is no bank conflict.

# CUDA

## Programming in CUDA

- Hardware-specific knowledge is required.
- Algorithms must be re-written to fit in the architecture.
- The optimal configuration changes depending on the algorithms.



# Haskell + CUDA

What if Haskell could provide a high-level abstraction that

- takes care of those low-level details,
- lets you focus on *WHAT to program* rather than *HOW to program*, and
- provides you with a similar performance?

# Haskell + CUDA

What if Haskell could provide a high-level abstraction that

- takes care of those low-level details,
- lets you focus on *WHAT to program* rather than *HOW to program*, and
- provides you with a similar performance?

GPGPU community

GPU programming becomes easier!

Haskell community

Free performance gain!

## GpuGen

How can we combine Haskell and CUDA?


# GpuGen

How can we combine Haskell and CUDA? **GpuGen**

# Overview

## GpuGen

- A compiler project in progress to compile Haskell to CUDA.
- Takes the external representation of Haskell's intermediate language, *Core*, as the input, and generates the corresponding optimised CUDA code.
- Will be plugged into *Nested Data Parallelism* (NDP) framework in *Glasgow Haskell Compiler* (GHC).



# Overview (Cont'd)

Which subset of the language is to be compiled?

The collective list operations are to be compiled:

- Most loops can be expressed with one or more collective list operations.
- The collective list operations well-fit into *Single Instruction, Multiple Data* (SIMD) and *Multiple Instruction stream, Multiple Data* (MIMD) models.
- Unsegmented and segmented  
map, scan, fold, filter, zipWith, permute, bpermute,  
and many more.


# Overview (Cont'd)

## Benefits

- Expression-level kernel extraction, resulting in functional program-decomposition (not hardware-based)
- Isolation of the side-effects
- Nested data processing (Flattening and Fusion)  
[[1, 5, 2, 3, 5], [], [13]]
- *What*, not *How*

# Overview (Cont'd)

## Design



1

<sup>1</sup>GpuExp is a *Generalised Algebraic Data Type* (GADT) to describe the operations on GPU.



# GpuExp

## GpuExp

- a GADT to describe what operations are to run on GPU.
- an *Embedded Domain Specific Language* (eDSL).

# GpuExp

## GpuExp

- a GADT to describe what operations are to run on GPU.
- an *Embedded Domain Specific Language* (eDSL).

## Saxpy in plain Haskell

```
xs, ys :: [Float] -- type signature
zipWith (\x y -> 5 * x + y) xs ys
```

## Saxpy in GpuExp

```
xs, ys :: Arr CInt (ForeignPtr CFloat)
run (ZipWith (\x y -> 5 * x + y) xs ys)
```

# GpuExp (Cont'd)

## GpuGen-generated Saxpy CUDA code

```
__device__ float _zipWithOpHSaxpy(float x, float y)
{
    return 5 * x + y;
}
```

# GpuExp (Cont'd)

## GpuGen-generated Saxpy CUDA code

```
__global__ void _zipWithHSaxpy
(const float *d_in0, const float *d_in1, float *d_out, uint n, uint isFullBlock)
{
    ...

    __syncthreads();
    i = aiDev * 4;
    tempData0 = inData0[aiDev];
    tempData1 = inData1[aiDev];
    if (isFullBlock || i + 3 < n) {
        tempData.x = _zipWithOpHSaxpy(tempData0.x, tempData1.x);
        tempData.y = _zipWithOpHSaxpy(tempData0.y, tempData1.y);
        tempData.z = _zipWithOpHSaxpy(tempData0.z, tempData1.z);
        tempData.w = _zipWithOpHSaxpy(tempData0.w, tempData1.w);
        outData[aiDev] = tempData;
    } else {
        if (i < n) {
            d_out[i] = _zipWithOpHSaxpy(tempData0.x, tempData1.x);
            if (i + 1 < n) {
                d_out[i + 1] = _zipWithOpHSaxpy(tempData0.y, tempData1.y);
                if (i + 2 < n) {
                    d_out[i + 2] = _zipWithOpHSaxpy(tempData0.z, tempData1.z);
                }
            }
        }
    }
    ...
}
```

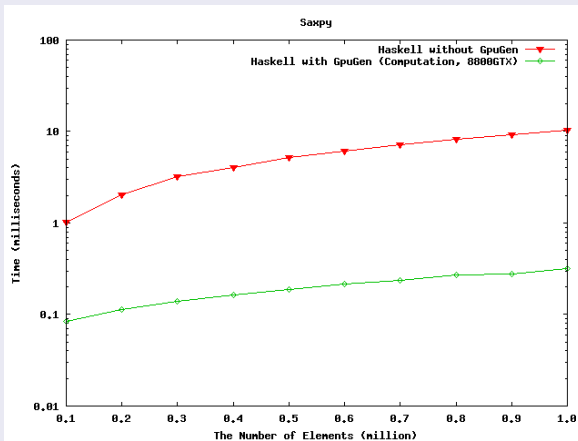
# GpuExp (Cont'd)

## GpuGen-generated Saxpy CUDA code

```
void gpuZipWithHSaxpy(float *xs0, float *xs1, float *res, uint n)
{
    ...
    CUDA_SAFE_CALL(cudaMalloc((void **) &d_xs0, memSize0));
    CUDA_SAFE_CALL(cudaMalloc((void **) &d_xs1, memSize1));
    CUDA_SAFE_CALL(cudaMalloc((void **) &d_res, memSize));
    CUDA_SAFE_CALL(cudaMemcpy(d_xs0, xs0, memSize0, cudaMemcpyHostToDevice));
    CUDA_SAFE_CALL(cudaMemcpy(d_xs1, xs1, memSize1, cudaMemcpyHostToDevice));
    CUT_CHECK_ERROR("Kernel execution failed");
    numBlocks = max(1, (uint) ceil((double) n / 512.0));
    _zipWithHSaxpy<<<dim3 (numBlocks), dim3 (64)>>>(d_xs0, d_xs1, d_res, n, n == numBlocks * 512);
    CUT_CHECK_ERROR("Kernel execution failed");
    CUDA_SAFE_CALL(cudaMemcpy(res, d_res, memSize, cudaMemcpyDeviceToHost));
    CUDA_SAFE_CALL(cudaFree(d_xs0));
    CUDA_SAFE_CALL(cudaFree(d_xs1));
    CUDA_SAFE_CALL(cudaFree(d_res));
}
```

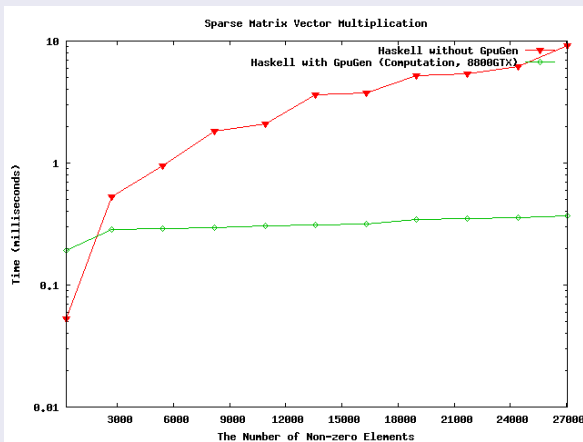
# Examples

## Saxpy



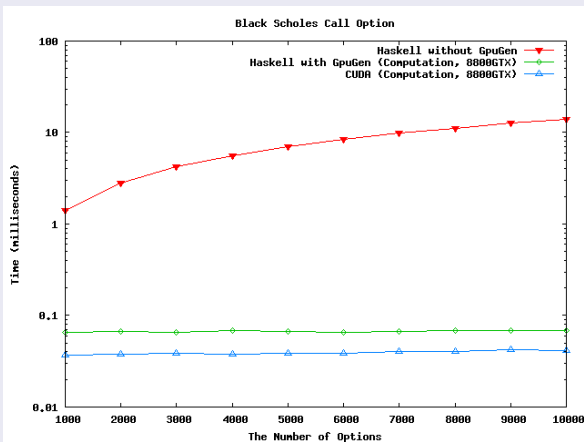
# Examples (Cont'd)

## Sparse Matrix Vector Multiplication



# Examples (Cont'd)

## Black Scholes Call Option





# Current Status

## Towards the first release

- GpuExp is in progress.
- Core parser is done.
- Translation of Core to CUDA is in progress.
- CUDA code generation for the collective list operations is done. (More operations could be added at any time.)

# GpuGen

## Fusion

### Minimises

- the data transfer between the host memory and the device memory, and
- the number of loops.

e.g. `map + scanl = scanl`

# GpuGen

## Fusion

### Minimises

- the data transfer between the host memory and the device memory, and
- the number of loops.

e.g. `map + scanl = scanl`

## Recursion

- Function pointers?

# GpuGen

## Fusion

### Minimises

- the data transfer between the host memory and the device memory, and
- the number of loops.

e.g. `map + scanl = scanl`

## Recursion

- Function pointers?

## CodeGen to PTX

- Haskell Core  $\rightarrow$  PTX
- More room for optimisations